

Lab 0

Deadline: Monday, June 29 (extended!!!)

Hello! Welcome to CS61C! We're excited to have you on board :D

This lab may be a little long, but please read carefully; it will prepare you for the rest of the course!

Objectives

Each lab will begin with a few objectives; use them to reflect on your learning and guide your thinking! Here are the objectives for this lab:

- The student will be able to describe and adhere to all course policies, including the lab checkoff policy.
- The student will set up accounts for GitHub, Gradescope, Hive machines, Piazza, and other course-related services.
- The student will gain familiarity with Linux and Git commands.

Jargon

A quick clarification on some terms:

- TA/GSI/uGSI: Teaching Assistant (sometimes called Graduate/Undergraduate Student Instructor).
- AI: Academic Intern, also part of course staff. You'll see them in checkoff, OH, and Piazza.
 - In this course, AI generally stands for this, and not "artificial intelligence".
- OH: Office Hours, where you can meet course staff in (virtual) meetings and ask 61C-related questions. Or occasionally, if it's not busy, non-61C-related things.
- Hive/"Hive machines"/"the Hive": a group of instructional servers. More details later in the lab.

Exercise 1: Course Policies

You can view the course policies on the [policies page of the course website](#). You are responsible for understanding and adhering to all policies throughout the course of the semester. Please pay particular attention to the [Grading](#), [Labs](#), and [Academic Dishonesty and Cheating](#) sections.

We have modified our lab checkoff policy; please read [New Lab Checkoff Policy](#) on Piazza for details.

Action Items

- Once you've finished reviewing the policies, take the [Lab00: Course Policies Quiz](#) on Gradescope.

Exercise 2: Setup

Unfortunately, assignments in this course do require some (sometimes boring) setup. Let's get that out of the way.

Accessing Services

CS 61C primarily uses a couple services, which are detailed below.

Unfortunately, some services and resources may be inaccessible or blocked in certain regions. As I'm writing this lab:

[Objectives](#)

[Jargon](#)

[Exercise 1: Course Policies](#)

[Exercise 2: Setup](#)

[Exercise 3: Command Line Essentials](#)

[Exercise 4: Fun with Git](#)

[Info: Lab Submission and Checkoff](#)

[Checklist](#)

[Appendix](#)

- GitHub (AWS + their own infrastructure) is inaccessible in some regions.
- Google and YouTube are inaccessible in some regions.
- Gradescope and Piazza are hosted on Amazon Web Services (AWS), which may be partially inaccessible in some regions.
- OH Queue and OkPy are hosted on Google Cloud Platform (GCP), which may be partially inaccessible in some regions.
- The 61C website is hosted by GitHub Pages with Cloudflare in the middle, so it may be partially inaccessible in some regions.
- Hive machines are accessed using the SSH protocol, which may be partially blocked in some environments (e.g. CalVisitor WiFi on campus).



If you're unable to access any services or resources, please download and use the Berkeley campus VPN at <https://www.lib.berkeley.edu/using-the-libraries/vpn>. If the campus VPN is inaccessible or doesn't work for you, contact us (Piazza, or email [cs61c \[at\] berkeley.edu](mailto:cs61c@berkeley.edu)) and we can try to work something out.

CalNet ID (and Berkeley Google accounts)

Most students should have a CalNet ID (and therefore, email ending in @berkeley.edu).

- If a service allows CalNet ID login, use that whenever possible.
- If a service prompts you to sign in with Google, log in using your @berkeley.edu email if possible. This is your "Berkeley Google account" (or bConnected account, but nobody says that).

If you would like to use an existing account on a service, try to add your Berkeley email as a secondary email if possible.

Piazza

[Piazza](#) is a discussion forum that we'll be using as the main method of communication for this course. All announcements will be made here, and almost all questions or comments you may have should be posted here (unless we say otherwise).

We've already emailed invites to the 61C Piazza course to most students. If you're not enrolled on our Piazza course, please refer to the policies section linked in the next paragraph.

Please read through the [Discussion Forum section of our policies](#). Please remember and follow the Piazza Etiquette!

Instructional Accounts and Servers (the Hive!)

The "Hive" is a group of servers maintained by the EECS department. Campus is not very open right now, but they live in Soda 330 if you ever want to give them a hug! You can interact with these servers remotely through SSH and the command line (more on that later). We've already installed most of the software we'll be using on the Hive machines, so you can work on most assignments on the Hive if you want! You can find a list of Hive machines at [Hivemind](#) (the names starting with [hive](#)). Every Hive machine shares the same network mount, so files will be synced across all the machines.

You will need to [sign up for an EECS instructional account](#), which you'll use to access the Hive. We'll also be using your instructional account username as your primary ID for checkoff (though you can use your @berkeley.edu email address as a backup).

Note: If you can't create an account for whatever reason, don't worry! Fill out the form in the [Non-standard Enrollment section](#), and continue without an account for now.

Gradescope

[Gradescope](#) is the platform we use for submitting and grading assignments. Programming assignments, homework, and quizzes will be submitted here. Exam results will also be returned here.

We've already invited most students to the 61C Gradescope course. To sign in, visit [Gradescope](#), click [Log In](#), then [School Credentials](#), then [CalNet ID](#).

Warning: Please ensure you set your [@berkeley.edu](#) email as your primary email. If not, we might not be able to find your Gradescope account, and you may receive an F in the class!

GitHub

[GitHub](#) is a hosted Git service we use for code distribution, along with GitHub Classroom tools.

If you have an existing GitHub account, feel free to use that; repositories created in this course are private, and anything you do for this course shouldn't affect the rest of your GitHub account.

If you don't have a GitHub account or want to make a new one, [sign up for a new account](#).

YouTube

Most lectures and other video resources will be uploaded to [YouTube](#). You will need to be signed into YouTube using your Berkeley Google account to view some videos. If you're unable to view a video, make sure you're using your Berkeley Google account.

Zoom

We will be holding meetings (OH, lab checkoff, project parties, lectures, discussions, etc.) over [Zoom](#). When signing in, use the [Sign in with SSO](#) option, enter [berkeley.zoom.us](#), and sign in with your CalNet ID.

OH Queue

Office Hours (OH), checkoffs, and project parties will be scheduled and managed through the [OH Queue](#). Please use your Berkeley Google account when logging in; non-Berkeley Google accounts will not be able to use the OH Queue properly.

Please read the blue info box on the OH Queue page.

Non-standard Enrollment

If you're enrolled in the course as a Berkeley student, you can ignore this blurb.

If you're a concurrent enrollment student, or otherwise not enrolled in CS 61C, please fill out [this form](#) so you can get access to our resources.

Action Items

- Sign up for and log into all the services described above.
- If you're getting tired of reading, try taking a short break :)

Exercise 3: Command Line Essentials

If you took CS61A and CS61B, you likely have some experience with a command line interface (CLI) and terminal commands, but we'd like to list some less common commands here that may come in handy during the course. For a review of the basic UNIX commands, [look over this](#)

[guide](#). Be sure to read and understand section B of the guide as well as the commands below, since you'll need them for Exercise 4.

Example commands will be formatted like:

```
$ echo "Hello world"
```

Typing that (without the `$`) in your terminal will run the command. In this case, it just prints `"Hello world"`.

CLI Shortcuts

As a CLI refresher, when typing commands or file paths:

- `<tab>` will autocomplete the current term
- `<up arrow>` and `<down arrow>` will allow you to refill commands you've used previously without typing them again.
- `<ctrl> + a` will move the cursor to the beginning of the current line (helpful for fixing mistakes)
- `<ctrl> + e` will move the cursor to the end of the current line (helpful for fixing mistakes)
- `<ctrl> + r` will let you search through your recently used commands

Working with Files

`touch` will create a blank file with the file name you provided.

```
$ touch example.txt
```

This will create a file named `example.txt` with nothing inside.

If you'd like to create a file with contents already inside, you can use:

```
$ echo "Your contents here, inside double quotes" > example.txt
```

This will create a file with the name `example.txt` in your current directory. If the file already exists, it will be overwritten. The file will contain `Your contents here, inside double quotes` but without the double quotes. The `>` symbol takes one argument which modifies where data printed to stdout is piped. Here, we are redirecting them to a file named `example.txt`.

You can also use the `echo` command by itself, in which case it will print the string to the terminal (without creating a file in the process).

```
$ echo "El Psy Kongroo"
```

You can view the contents of a file with the `cat` command.

```
$ cat example.txt
```

This will print out the file contents of `example.txt` to your terminal. The file must be in your current directory, but you can provide a relative or absolute path to print out non-local files.

man - Manual Pages

The manual pages (“man pages”) are great UNIX resources that are often underused; while not as versatile as Google, they contain documentation on UNIX components from program usage, language standards and conventions, and more. In this course we’d like you to get comfortable with them, especially for C and UNIX-related questions.

If you want the man page for a single program/command, you can run:

```
$ man command_name | less
```

The man page for a program typically contains information about what the program is used for, what certain flags do when you invoke the program with them, and where to go for more information. Since we piped the man page into `less`, this page is scrollable (use your arrow keys or the space bar). Hit `q` to exit the man page and get back to your terminal prompt.

```
$ man echo | less
```

The above command should bring up the man page for the `echo` command.

If you want to search the man pages for a command that pertains to a keyword:

```
$ man -k single_keyword | less
```

This command will search the manual pages for a command with the keyword `single_keyword`. Forget how to open files in Vim? You can search for `editor` and get a list of all editor-related commands on your system.

ssh - “Secure Shell”

For this class, we’ll expect you to test most of your projects, homeworks, and labs on the Hive servers. To access the Hive servers remotely, you’ll be using the SSH protocol and programs.

Note: If you weren’t able to get an instructional account, you can come back here later!

You can find a list of Hive machines at [Hivemind](#). There are 30 of them, named `hive1`, `hive2`, ..., `hive30`. Sometimes, a Hive machine may be down or overloaded; when this happens, you can check Hivemind and find another machine to use.

Once you have an instructional account, you can SSH into an instructional server with the following command:

```
$ ssh cs61c-???@hive#.cs.berkeley.edu
```

Remember to replace `cs61c-???` with your instructional account username, and `hive#` with a Hive machine’s name. The default password is displayed by [WebAccount](#) when creating the account, so you might have to reset your password if you forgot it.

Troubleshooting:

- If nothing happens for a long time: check your internet connection. Some firewalls and networks, including [CalVisitor](#) on campus, block SSH. Try another network ([AirBears2](#) or [eduroam](#) if you’re on campus).
- `Connection refused` or other weird errors: the Hive machine you picked might be down. Try another one
- `Reserved for cs61c staff`: try another Hive machine :)

When your connection succeeds, you should be able to interact with and run commands on your chosen Hive machine! To exit this SSH session, simply run:

```
$ exit
```

Files on the Hive machines are stored on a network drive, so your account will have the same files on all 30 machines.

If you want to change your instructional account password, you can SSH into the update server:

```
$ ssh cs61c-???@update.cs.berkeley.edu
```

scp - “Secure Copy”

The `scp` program is used for copying files between computers using the SSH protocol.

Sometimes, you may want to get individual files or entire folders from the Hive machines onto your local system, or vice versa. You can do this by using `scp`:

```
$ scp <source> <destination>
```

To specify a remote source or destination, use `username@host:path`. To specify a local path, just use `path`. As an example:

```
$ scp cs61c-xyz@hive3.cs.berkeley.edu:~/some_folder/example.txt ~/Downloads/
```

Assuming my username is `cs61c-xyz`, the above command would connect to `hive3` and copy `~/some_folder/example.txt` on my instructional account to `~/Downloads/example.txt` on my local machine.

If I wanted to copy the other direction (from my local machine to a Hive machine) I would use:

```
$ scp ~/Downloads/example.txt cs61c-xyz@hive7.cs.berkeley.edu:~/some_folder/
```

Warning: you shouldn't run `scp` on the Hive machines (when you're in a SSH session). Always run it in a local terminal session!

`scp` by default only works with files. To copy folders, you need to tell `scp` to “recursively” copy the folder and all its contents, which you can do with the `-r` flag:

```
$ scp -r cs61c-xyz@hive7.cs.berkeley.edu:~/some_folder ~/Downloads/
```

Vim Basics

`vim` is a text editor included on the Hive machines and many UNIX-based distributions.

Note: We'll be using Vim in most of our examples and documentation, but we have no hard requirement on which text editor you use; you're welcome to pick whatever you're comfortable with, but you should know how to use at least one terminal-based text editor.

To open a file from your current directory, pass the file name to Vim:

```
$ vim filename
```

To open a file from another directory, use a relative or absolute path:

```
$ vim ../other_folder/filename
```

Some useful Vim commands:

Command	Explanation
<code><escape>:q</code>	Closes (quits) Vim without saving
<code><escape>:wq</code>	Closes Vim after saving
<code><escape>:w</code>	Saves your file
<code><escape>:q!</code>	Force-quit Vim (for when you've made changes but do not wish to save them)
<code><escape>i</code>	Insert mode, allows you to type into the file
<code><escape>/cats</code>	Searches your file for the nearest occurrence of the string "cats". Press <code>n</code> to go to the next occurrence or <code>N</code> to go to the previous
<code><escape>:set nu</code>	Shows line numbers within your file

Note: these commands are preceded by `<escape>` because you'll need to press the escape key on your keyboard to switch you out of your current mode. For example, if I'm inserting (typing) into a file and want to save, I'd have to hit `<escape>` to get out of insert mode, then type `:w` to save my file. If you aren't in a mode (i.e. you've just opened your file) you don't need to hit escape first, but it won't hurt :)

If you want to learn more about Vim, one of our tutors, Yijie, wrote a great [Vim for CS61C guide!](#)

Action Items

SSH into any Hive machine. Then:

- If there is a prompt asking you to enter some information:
 - Last name (family name)
 - First name (given name) and any middle name(s)
 - Student ID
 - Email address: please use your Berkeley email
 - Code name: just pick something random, we don't use this
- Whether or not there was a prompt, run `check-register`, and verify that your name, email, and student ID are correct. If anything is incorrect, run `re-register`.
- Make sure the email address you entered is listed as a primary or secondary email on your Gradescope account.

Exercise 4: Fun with Git

This exercise will show you how to set up a Git repository ("repo"), use Vim, and work with a variety of Git commands. By the end of it, you should feel comfortable using SSH, editing files, pulling/committing/pushing, resolving merge conflicts. If you'd like to review your Git commands before beginning, you can check out [this guide](#).

Getting Your Lab Repo

Make sure you're logged into GitHub. Fill out the [Lab Repository Registration Form](#). Remember to read the form carefully, and follow the link in the form to accept our invitation to the GitHub Classroom lab assignment. You may need to check your email and confirm the invitation. If you don't have an instructional account right now, enter `zzz` in the form. You will be able to edit this later.

Configuring Git

Before we start, let's tell Git who you are. This information will be used to sign and log your commits. You may not need to do this if you've set up Git before, but if you're on the Hive machines it's likely a step you'll need to take.

First, run the following commands on your local machine; make sure to change the name and email to match your information.

```
$ git config --global user.name "John Doe"
$ git config --global user.email johndoe@example.com
```

If you have an instructional account, SSH into a Hive machine, and run the same commands.

Cloning Your Repo

Git has the concept of “local” and “remote” repositories. A local repo is located wherever your terminal session is; if you're in a SSH session, the local repo is a folder on a Hive machine; if your terminal session on your local machine, the local repo is located on your local machine's filesystem. A remote repo (e.g. GitHub repo) is typically hosted on the Internet.

GitHub Classroom created a repository for you on GitHub, but not locally. To get a local copy of this repository, you can use `git clone`, which will create a local repository based on information from a remote repo.

If you have an instructional account, SSH into a Hive machine. On the Hive clone the repository into a folder called `labs`:

```
$ git clone <link to repo created by GH Classroom> labs
```

If you don't have an instructional account, that's fine! Clone the repository into a folder called `labs_hive`. For the rest of this exercise, any reference to your repository on the Hive is referring to this repository.

```
$ git clone <link to repo created by GH Classroom> labs_hive
```

Exploring Your Repo

`cd` into this new folder. List all hidden files (`ls -a`). Do you see a hidden file/folder?

There is indeed a folder called `.git`. Its presence indicates that the current folder (folder containing `.git`) holds a Git repository.

Take a look at your repo's current remotes and status:

```
$ git remote -v
$ git status
```

`git clone` has already worked a bit of magic here; there's a remote called `origin`, and its URL points to your labs repo on GitHub! You're on a local branch called `master`, which is “tracking” `origin/master` (the `master` branch on the `origin` remote).

Throughout the semester, course staff may make updates to starter code to fix bugs or release new labs. To receive these updates, you'll need to add another remote.


```
$ git remote add starter https://github.com/61c-teach/su20-lab-starter.git
```

If you ever want to pull **updated** starter code, you'd execute the following command:

```
$ git pull starter master
```

Since GitHub Classroom clones the starter repo when creating your repo, that command shouldn't do anything right now.

Finding Commands

Warning: please read carefully. Skipping a step may cause errors that will require you to redo the exercise.

The files for this exercise are located in the `lab00` folder in your lab repository.

1. Open `first_set.txt`. This file contains descriptions of UNIX commands. Use the man pages to figure out which command they refer to; **be sure to include any necessary flags**. Once you've found them, write the commands, one per line, in the provided `answers.txt` file.

`answers.txt` should look similar to:

1. `first_set_command_1`
2. `first_set_command_2`
3. `first_set_command_3`
4. `first_set_command_4`

Save, add, and commit your changes with the commit message "`Answered first set`". Push this commit to your GitHub repo.

2. Now, we'll need to copy the repository to your local computer.
 - If you have an instructional account, exit the SSH session by running `exit`. You should be back to your local terminal prompt. `scp -r` the entire `labs` folder from the Hive to your local machine. Remember why we need the `-r`!
 - If you don't have an instructional account, copy the entire `labs_hive` folder to a separate `labs` folder.

Once you've copied the entire folder, `cd` into this folder and open your answers file.

3. **Above** your previous answers, add commands (and relevant flags, if needed) that fit this second set of descriptions:
 - Move files from one directory to the next
 - Change a file's permissions
 - Show the directory you're currently in
 - Check if a host is online

`answers.txt` should look similar to:

1. `second_set_command_1`
2. `second_set_command_2`
3. `second_set_command_3`
4. `second_set_command_4`
5. `first_set_command_1`
6. `first_set_command_2`
7. `first_set_command_3`
8. `first_set_command_4`

Save your changes, but **do not commit or push them**.

4. Go back to your repo on the Hive. **Above** your previous answers, add the command (and relevant flags, if needed) that fits this description:

- Show file permissions for all files in a directory

`answers.txt` should look similar to:

1. `third_set_command_1`
2. `first_set_command_1`
3. `first_set_command_2`
4. `first_set_command_3`
5. `first_set_command_4`

Save, add, commit, and push these changes with the commit message "`Answered third set`".

5. Go back to your repo on your local machine. Add, commit, and push your changes with the commit message "`Answered second set`".

You should get an error message from Git. What does this message mean? How do you resolve it?

After resolving the issue, you should have all your answers in `answers.txt` on your local machine. `answers.txt` should match the following descriptions, with the numbers:

1. Show file permissions for all files in a directory
2. Move files from one directory to the next
3. Change a file's permissions
4. Show the directory you're currently in
5. Check if a host is online
6. Which command walks a file hierarchy in search of a keyword?
7. Which command displays information about processes running on your machine?
8. Which command terminates a process?
9. Which command can help you find the difference between two files?

Commit and push your changes.

6. Still on your local machine, from the `lab00` directory, run:

```
$ ./copy_git.sh
```

This will add, commit, and push a log of your Git repo. If this succeeds, you should be able to see a file called `git.log` in your GitHub repo.

Info: Lab Submission and Checkoff

Last bit of reading for this lab, I promise.

Submission

To submit your work, push your work to your lab repository on GitHub.

Many of the labs will have an autograder portion. If a lab has an autograder portion, it will specify this in the `Checklist` section at the bottom of the page. To submit to the autograder, you'll need to push your work to your lab repository on GitHub. Then go to the `Lab Autograder` assignment on Gradescope, and submit your lab repository. After a short wait, the page should show your autograder score for the current lab (and previous labs). Make sure you passed the autograder tests!

Checkoff: Overview

Each lab will require you to complete a checkoff to receive credit. ~~You will be required to work on labs in pairs, and both partners are required to be present during checkoff.~~ This helps us reduce the time we spend on checkoff and allocate that time towards helping students and developing course materials. It also helps you by giving you someone to discuss class material with and work together with to solve lab problems.

If you have not done so already, please read through the [Labs section of our policies](#).

We have modified our lab checkoff policy; please read [\[Important\] New Lab Checkoff Policy](#) on Piazza for details.

Checkoff: Requirements

Before your scheduled checkoff, please make sure that:

1. You've completed the action items for each exercise in the lab.
 - If the exercise asks any questions, make sure you've recorded your answers so you're prepared to answer these questions during checkoff.
 - If the exercise asks you to modify code, make sure you've modified code as necessary, and that you've pushed your modifications to your GitHub lab repo.
2. If the lab has an autograder, you've submitted your GitHub repo to the [Lab Autograder](#) assignment on [Gradescope](#). Your work should pass the autograder tests, and you should have this page open for checkoff.
3. You (and your partner, if you have one) are available for the checkoff slot you've selected.
4. You know your email address and instructional account username (`cs61c-???`, the one you use to log into Hive machines).

Checkoff: Signing Up

We will be using an appointment-based system on the OH Queue to schedule checkoffs. Checkoff slots for the week will be released on Sunday around 9PM PDT. You (and a partner) may then sign up for a 13 minute checkoff slot. If you have a partner, you both **must** sign up for the same slot (sometimes there are multiple slots at the same time, watch out!). You **must** have completed the requirements before your checkoff slot starts.

To sign up for a checkoff slot:

- Visit the [OH Queue](#) and click the [Appointments](#) link at the top.
- Find **one** slot that's suitable for you, then click [Add yourself to the section](#). Fill out the resulting form (Assignment: [Lab ##](#), Question: [Checkoff](#), Description: [Partner: <PARTNER_NAME>](#)) and submit it. If you have a partner, you both **must** sign up for the same appointment (try to sign up at the same time). If you don't have a partner, you may be paired with someone else without one. You may not sign up for more than one appointment across Tuesday+Wednesday, and one appointment across Thursday+Friday.
- When it's time for your appointment, go to the OH Queue. Your appointment will show up under [Upcoming Appointments](#) on the home page. When the TA is ready, a green [Join Call](#) button will appear. Click it to be taken to the Zoom call.

Checkoff: Receiving Credit

You will receive credit for a lab after completing a checkoff with a TA/Al. You can (and should) verify this by looking at the [Lab## Checkoff](#) assignment on Gradescope, replacing `##` with the lab number, and making sure it is [Submitted](#). It is your responsibility to verify that this happens for each lab; we will not be retroactively giving credit for checkoff.

If you completed checkoff in the correct timeslot, but Gradescope displays a [Submitted ????](#) [late](#) message, you can ignore this message for now; we'll update this after the lab deadline has passed.

Checklist

You made it! That was quite a bit of reading and head-scratching, but you're now somewhat more familiar with the tools you'll be using for the rest of the semester. Worth it!

This lab has an autograder. After you've pushed your work, submit your GitHub lab repo to the [Lab Autograder](#) assignment on Gradescope. Make sure the autograder passes!

Please check that you and your partner have:

- Completed the course policies quiz from Exercise 1.
- Registered for the services from Exercise 2.
- Registered your information on the Hive in Exercise 3.
- Completed Exercise 4 (matching the described at the end).
- Passed the [Lab Autograder](#) for this lab

After a TA/Al has checked you off, please make sure that the corresponding lab assignment on Gradescope ([Lab## Checkoff](#)) is marked as [Submitted](#).

Appendix

These are some tools you may find helpful, but are by no means required for this course :)

Text Editor vs. IDE

CS61C doesn't endorse any particular text editor or IDE. Many people get by in this course using a text editor with no frills (think: Vim/Emacs/Nano). We'll expect you to know how to use at least one terminal-based text editor (again, Vim/Emacs/Nano), since you'll be dealing with the command line a lot.

For your own work, you may find it nice to have CLion from JetBrains if you're used to working in IntelliJ from CS61B. Note though that we won't be providing any course-official support, so setting it up and maintaining it are up to you.

The majority of students do their work in a local editor (Sublime, Atom, VSCode) and use Git to copy their files from their local machine to the Hive. Some students also set up Cyberduck to make copying files over easier. Again, we won't provide any course-official support, but you're welcome to do what works best for you.

Quick SSH

Tired of typing up an entire SSH command and password? Follow the instructions on [this Piazza post](#) ("Useful additional setup information for labs and projects").